

HPX Backend for Blaze

Shahrzad Shirzad

Nov 7, 2019

Linear Algebra Library based on Smart Expression Templates

- Expression Templates:
 - Creates a parse tree of the expression at compile time and postpone the actual evaluation to when the expression is assigned to a target
- Smart:
 - Creation of intermediate temporaries when needed
 - Integration with highly optimized compute kernels

Depending on the operation and the size of operands, the assignment could be parallelized through four different backends

- HPX
- OpenMP
- C++ threads
- Boost

In the current implementation the work is equally divided between the cores at compile time.

- HPX for-loop with static chunking and chunk size=1




Figure 1: An example of how $C=A+B$ is performed in HPX Backend with 4 cores

Objective

Dynamically divide the work among the cores based on number of cores, matrix size, operation, etc. For this purpose two parameters have been introduced:

- `block_size`: at each loop iteration the assignment is performed on one block
- `chunk_size`: the number of loop iterations included in one task




Figure 2: An example of blocking and creating chunks for `chunk_size = 2`

- Effect of Task Granularity on execution time
- Universal Scalability Law

Task Granularity

Grain size: The amount of work performed by one HPX thread

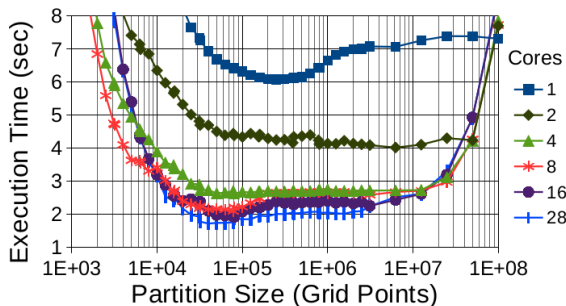


Figure 3: The effect of task size on execution time for Stencil application¹

¹Grubel, Patricia, et al. "The performance implication of task size for applications on the hpx runtime system." 2015 IEEE International Conference on Cluster Computing. IEEE, 2015.

Universal Scalability Law

- Models the effects of linear speedup, contention delay, and coherency delay due to crosstalk

$$X(N) = \frac{\lambda N}{1 + \sigma(N-1) + \kappa N(N-1)}$$

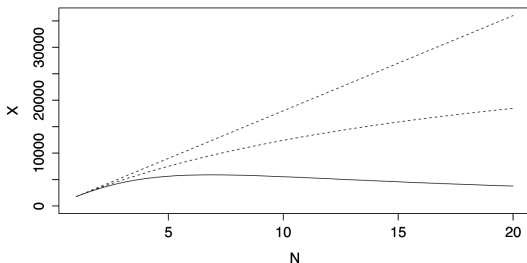


Figure 4: Throughput vs. number of cores²

²Schwarz, B. "Practical Scalability Analysis with the Universal Scalability Law." (2015).

Blazemark

Blazemark is a benchmark suite provided by Blaze to compare the performance of Blaze with other linear algebra libraries.

```
Dense Vector/Dense Vector Addition:
C-like implementation [MFlop/s]:
  100      1115.44
 10000000  206.317
Classic operator overloading [MFlop/s]:
  100      415.703
 10000000  112.557
Blaze [MFlop/s]:
  100      2602.56
 10000000  292.569
Boost uBLAS [MFlop/s]:
  100      1056.75
 10000000  208.639
Blitz++ [MFlop/s]:
  100      1011.1
 10000000  207.855
GMM++ [MFlop/s]:
  100      1115.42
 10000000  207.699
Armadillo [MFlop/s]:
  100      1095.86
 10000000  208.658
MTL [MFlop/s]:
  100      1018.47
 10000000  209.065
Eigen [MFlop/s]:
  100      2173.48
 10000000  209.899
```

```
N=100, steps=55116257
C-like      = 2.33322 (4.94123)
Classic     = 6.26062 (13.2586)
Blaze       = 1      (2.11777)
Boost uBLAS = 2.4628 (5.21565)
Blitz++     = 2.57398 (5.4511)
GMM++       = 2.33325 (4.94129)
Armadillo   = 2.3749 (5.0295)
MTL         = 2.55537 (5.41168)
Eigen       = 1.19742 (2.53585)
N=10000000, steps=8
C-like      = 1.41805 (0.387753)
Classic     = 2.5993 (0.710753)
Blaze       = 1      (0.27344)
Boost uBLAS = 1.40227 (0.383437)
Blitz++     = 1.40756 (0.384884)
GMM++       = 1.40862 (0.385172)
Armadillo   = 1.40215 (0.383403)
MTL         = 1.39941 (0.382656)
Eigen       = 1.39386 (0.381136)
```

Figure 5: An example of results obtained from Blazemark

- Starting from *DMATDMATADD* benchmark, $C = A + B$ where A , B , and C are N by N matrices
- Collect data with different configurations such as matrix size, number of cores, block_size, chunk_size.
 - matrix sizes: 200, 230, 264, 300, 396, 455, 523, 600, 690, 793, 912, 1048, 1200, 1380, 1587
 - number of cores: 1, 2, 3, ..., 8
 - block_sizes: [4, 8, 12, 16, 20, 32] by [64, 128, 256, 512, 1024] blocks
 - chunk_sizes: between 1 and total number of blocks (logarithmic increase)

Results

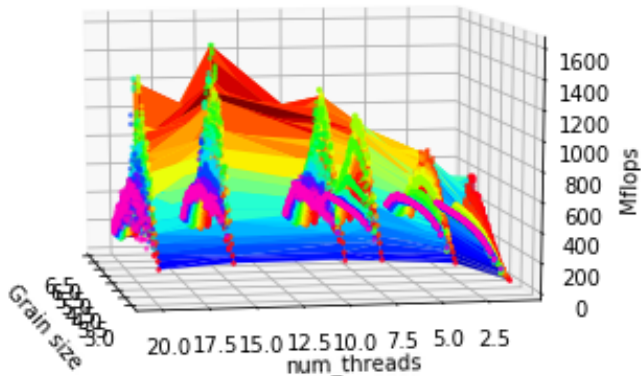


Figure 6: Results of running the *DMATDMATADD* benchmark for different matrix sizes with different *block_size* and *chunk_size* combinations

Method

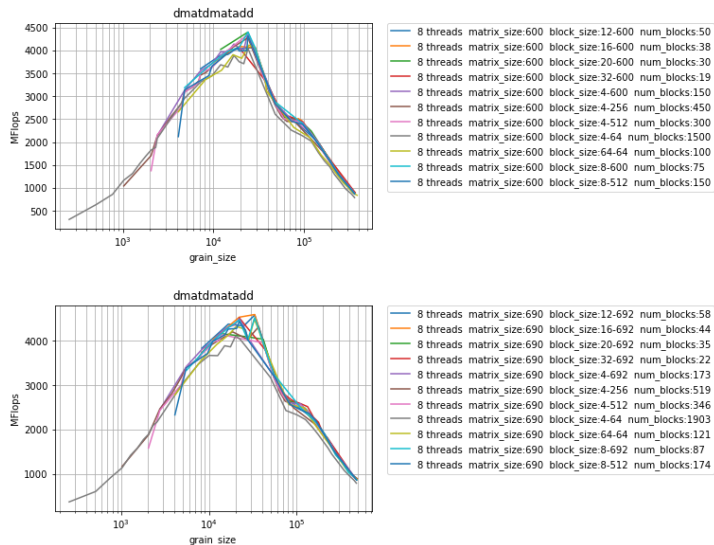



Figure 7: Results of running the *DMATDMATADD* benchmark on 8 cores for different block sizes

Modeling Execution Time based on Grain Size



Modeling Execution Time based on Grain Size

- Overheads of creating tasks
- Starvation

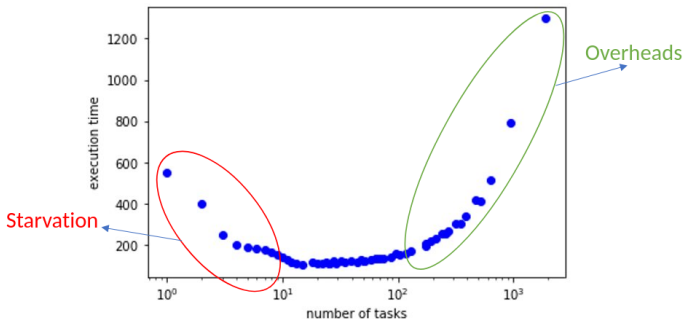


Figure 9: Results of running the *DMATDMATADD* benchmark on 8 cores matrix size 690 (time unit is microseconds)

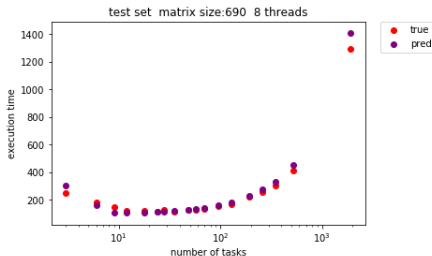
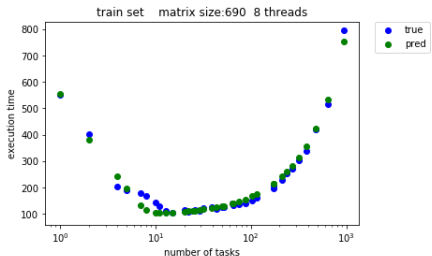
Modeling Execution Time based on Grain Size

Execution_time =

$$\begin{cases} \frac{\alpha \times num_tasks + t_s}{num_tasks} + \beta \times num_tasks + \gamma & num_tasks < N \\ \frac{\alpha \times num_tasks + t_s}{N} + \beta \times num_tasks + \gamma & num_tasks \geq N \end{cases}$$

Modeling Execution Time based on Grain Size

- Fixed matrix size, and number of cores
- Training set and test set (%60, %40)



Modeling Execution Time based on Grain Size

- For a fixed matrix size, and number of cores we need 4 parameters to estimate execution time based on number of tasks
- How does these four parameters change for different number of cores?
 - used USL to model each of these parameters

Modeling Execution Time based on Grain Size

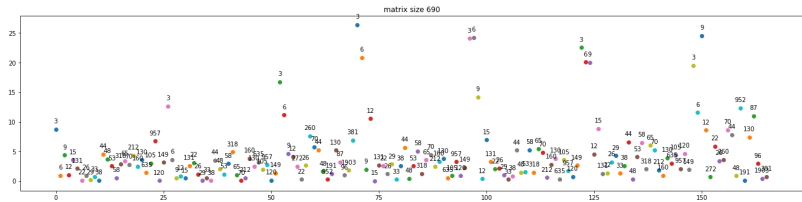


Figure 11: Relative error of predicting execution time

Next Step

- Find the range of the flat region of grain size
- Choosing a small `block_size` while number of columns is divisible by cache line
- Find the range of `chunk_sizes` for the given range of grain size
- Generalize the model to integrate the matrix size

Thank you!